

TorchFDTD: GPU finite-difference time-domain simulation with discrete adjoints and tiered-memory execution for photonic design

Hyoseok Park

Department of Physics, Chungnam National University, Daejeon, Republic of Korea

21 September 2026

Abstract

Gradient-based photonic design needs a time-domain solver that is fast on one graphics processor, differentiable with respect to the design, and able to run problems whose field state exceeds the memory of that processor. We present TorchFDTD, an open-source finite-difference time-domain (FDTD) package that couples a native CUDA Yee engine to PyTorch. The engine advances the six field components with fused kernels captured in CUDA Graphs and supports convolutional perfectly matched layers with independent per-face profiles, periodic and Bloch boundaries, perfect electric and magnetic conductor faces, multipole Drude and Lorentz dispersion, anisotropic permittivity tensors admitted inside the absorbing layers under a geometric stability criterion, graded rectilinear meshes, subpixel interfaces, and total-field/scattered-field, one-way and guided-mode sources. First-order gradients follow from an explicit transpose of every discrete update combined with binomial checkpoint replay, and they cover material, density, polygon and spline shape, and source-waveform parameters together with objectives on point signals, plane spectra, multiport scattering parameters and projected far fields. A streamed execution mode partitions the domain into slabs, keeps only the active slabs on the GPU, banks the remaining state in host memory or on solid-state storage under measured reservations, and writes a durable journal from which an interrupted run resumes. On an RTX 5880 the fused engine solves eight forward fixtures 9.6 to 17.1 times faster than a CUDA-graph adaptation of an existing PyTorch FDTD package, and a 256^3 adjoint runs 179 times faster than the same Torch model on twelve CPU threads. On one RTX 3060, forward solves run 6.4 to 7.3 times faster than FDTDx and 38 to 44 times faster than twelve-rank Meep. A 2.26 billion cell single-precision adjoint whose 54 GB of field state exceeds the 51.5 GB of the processor completes with a 3.0 GB peak device allocation, and a process killed after its first backward record resumes from the journal with a gradient within 9×10^{-8} of the uninterrupted oracle. Analytic slab spectra, Mie scattering, dipole radiation, waveguide scattering parameters and finite-difference gradient checks bound the numerical error of each path.

1 Introduction

The finite-difference time-domain (FDTD) method advances Maxwell’s equations on a staggered grid and returns broadband observables from one run [1, 2]. Absorbing layers [3, 4], auxiliary differential equations for dispersive media [5] and subpixel material averaging [6] made the method the standard forward solver of nanophotonics, and mature packages such as Meep [7] and commercial tools [8, 9] expose it through scripting interfaces. Inverse design shifted the requirement from a forward solver to a differentiable one: adjoint sensitivities turn a single extra simulation into a gradient over millions of design variables [10–14]. The Fourier modal solver TORCWA showed how a PyTorch implementation makes such gradients available to ordinary optimizers on a graphics processor [15], and FDTDx brought reverse-mode differentiation and multi-GPU sharding to time-domain simulation in JAX [16].

Two costs remain with the time-domain adjoint. The first is memory. The backward pass of an FDTD run needs the forward field history, and a time-reversible reconstruction or a checkpointing schedule trades that history against recomputation [16, 17]. Either way the peak state of a three-dimensional problem,

six field components plus absorber and material memories over every cell, must fit the processor, so the largest design problem is fixed by the size of one graphics card. The second cost is the granularity of the differentiation itself. A solver written entirely in an automatic-differentiation framework differentiates every operation it executes, so a fused hand-written CUDA kernel that removes the memory traffic of the framework also removes its gradient unless the transpose of that kernel is written by hand.

TorchFDTD addresses both costs with one engine. The Yee, absorber and dispersion updates are fused CUDA kernels whose transposes are written explicitly, so the adjoint has the same kernel-level efficiency as the forward pass while remaining a first-class PyTorch operation. The field state is partitioned into slabs along one axis, and a streamed execution mode keeps only the slabs under update on the device, banks the rest in host memory or on solid-state storage, and replays checkpoints at the block level. The reservations of that hierarchy are measured rather than assumed: the number of simultaneously live file banks is bounded by the checkpoint count plus three, and the host-side tensors of the dense material gradient are counted from a recorded allocation ledger. A journal of block-boundary records makes a multi-hour adjoint resumable after a process failure. Around the engine sit a serializable scene model shared by a browser workbench and a Python interface, ensembles that execute independent structures in one CUDA launch, guided-mode ports with multiport scattering parameters, near-to-far field projection, GDSII import and export, and an independent reader and writer for a documented subset of a commercial project format.

This paper describes the discrete formulation (Section 2), the execution and differentiation machinery (Sections 3 and 4), the tiered-memory mode (Section 5), the workbench (Section 6), and the validation and performance records (Sections 7 and 8). Every number quoted below is taken from a JSON record in the repository, which also stores the source hashes of the code that produced it [18]. The grid and backend foundation come from the MIT-licensed `flaport/fdtd` package [19], and PyTorch supplies device arrays, graph capture and the optimizer interface [20].

2 Discrete formulation

2.1 Yee update and time step

The six components occupy the standard staggered positions [1]. In a nonmagnetic medium the reduced update reads

$$E^{n+1} = E^n + S \epsilon^{-1} C^- \left(H^{n+1/2} \right) + s^{n+1}, \quad H^{n+3/2} = H^{n+1/2} - S C^+ \left(E^{n+1} \right) + m^{n+3/2}, \quad (1)$$

where $S = c\Delta t/h$ with h the smallest mesh step, $C^\pm$ are the forward and backward difference curls scaled by the local metric, and s and m are soft electric and magnetic source increments added after the corresponding update. Fields start from zero, the first recorded electric sample lies at Δt and the first magnetic sample at $3\Delta t/2$, and every Fourier transform in the package keeps that half-step offset. For d active dimensions with minimum primal widths h_a the time step is bounded by

$$\Delta t \leq \frac{\eta}{c \sqrt{\sum_{a=1}^d h_a^{-2}}}, \quad 0 < \eta \leq 0.99, \quad (2)$$

which reduces to $\eta h/(c\sqrt{d})$ on a uniform mesh. A graded rectilinear mesh keeps this conservative fine-grid step, so coarsening the background reduces the cell count at fixed physical duration without local time stepping. The two-dimensional solver suppresses the z derivatives while keeping all six components, and Bloch problems carry complex fields.

2.2 Boundaries

Convolutional perfectly matched layers [4] store one memory variable per transverse derivative inside each absorbing face. The stretching coefficients κ , σ and α follow polynomial profiles that are configured per

face, so opposite faces and neighbouring faces of different type may carry different depths and gradings. Periodic and Bloch pairs wrap the difference operators without a duplicated end plane, and a Bloch phase enters through the wrapped stencil. Perfect electric conductor faces and electric antisymmetry planes hold the tangential electric field at zero on the boundary node. Perfect magnetic conductor and magnetic symmetry faces require the dual: the upper boundary node of each electric component tangential to the face is a stored degree of freedom, and the package carries those face arrays through the resident, streamed and batched paths, through their checkpoints and journals, and through the absorber memories of neighbouring faces. Because any face type may adjoin any other, the reflection of an absorbing face next to a magnetic wall is measured rather than inferred (Section 7).

2.3 Materials

A nondispersive cell carries a real permittivity at each Yee electric position. Dispersive materials add up to sixteen Drude or Lorentz poles integrated with the trapezoidal auxiliary differential equation of Okoniewski et al. [5], and measured refractive-index tables are fitted to a passive multipole model whose error and band are reported before the fit is accepted. Anisotropic media use a symmetric positive-definite permittivity tensor sampled at grid nodes and applied to the curl through a node assembly $S = D^{-1/2}(\frac{1}{8} \sum R^\dagger \epsilon^{-1} R) D^{-1/2}$, whose eigenvalues stay at or above one so that the vacuum time step of Eq. (2) remains admissible. Inside an absorbing layer the stretched-coordinate formulation is unstable for tensors whose principal axes are rotated relative to the face normal or whose normal eigenvalue is the intermediate one, a consequence of the backward-wave criterion of Bécache et al. [21]. TorchFDTD therefore admits a tensor inside a face only when the face normal is a principal axis with a non-intermediate eigenvalue, a rule checked node by node. Dense one-step spectral radii on small Bloch grids equal 1.000000000 for admitted tensors and 1.0049 to 1.0089 for rejected ones, and a 4000-step sweep over fifteen admitted cases with contrasts up to sixteen shows no growth while the rejected cases grow by factors of 7.9×10^3 and 7.6×10^{12} . Tensor dispersion couples a symmetric positive semidefinite strength tensor to each pole, and poles inside an absorbing face require an axis-aligned ϵ_∞ with strengths proportional to it.

Geometry enters through analytic solids (boxes, cylinders, spheres, ellipsoids, rings and extruded polygons with holes) placed by an ordered rotation and a pivot, with a lower mesh order winning overlaps. Membership is evaluated only inside a conservative axis-aligned support, and an optional dielectric subpixel operator averages the permittivity across interfaces [6]. GDSII layouts map layer and datatype pairs to extrusion intervals and materials, apply Boolean etch layers, staircase a sidewall angle on the native z nodes, and keep polygon holes with even-odd filling [22].

2.4 Sources, monitors and projections

Point and sheet sources inject Gaussian, ramped continuous or sampled waveforms with Cartesian or angular polarization. A one-way plane covering a periodic transverse cell injects a normally incident wave without a backward component, and a closed total-field/scattered-field box surrounds an isolated scatterer with an incident field that is corrected on its faces. Guided-mode sources solve the full-vector finite-difference eigenproblem of the port cross-section [23] at the temporal wavenumber $\tilde{k}_0 = 2 \sin(\omega \Delta t / 2) / (c \Delta t)$, convert the eigenvalue to the longitudinal Yee propagation constant $\beta = 2 \arcsin(\tilde{\beta} \Delta w / 2) / \Delta w$, and inject electric and magnetic Huygens sheets at their staggered positions and times. The eigensolver returns a canonical basis for degenerate polarization pairs by diagonalizing the overlap of their transverse components, so a port's mode index has the same meaning on every machine.

Point monitors record one component per step. Plane monitors accumulate six-component complex spectra with the Yee interpolation and the magnetic half-step phase applied before the signed Poynting flux is formed, and flux ratios are taken against a matching reference run. A closed six-face box of stored spectra feeds a near-to-far transform of the equivalent surface currents in the asymptotic form and in a finite-distance form built on the free-space dyadic Green function, with angular, spherical, Cartesian and direction-cosine observation sets, a complex passive exterior index that may vary per frequency, and an

open-surface variant that projects from one to five faces with an edge window. Periodic cells return Bloch diffraction orders separated into forward and backward propagating branches.

3 Execution

One simulation on the device runs as a sequence of fused kernels: one kernel updates all electric components together with the absorber memories of every face, and one kernel updates the magnetic components. The kernels are captured once into a CUDA Graph together with source injection and monitor accumulation, so a run of N steps replays N graph launches without Python in the loop [20]. The CPU path executes the same discrete operators with NumPy and Torch tensors and serves as the reference for every device test.

Design workflows evaluate many structures, so the package provides two levels of concurrency. A process-level batch runner executes independent projects in separate processes on one or more devices with checksum-validated resumption, per-case error capture and a memory estimate that limits the number of resident cases. A tensor batch instead places B compatible structures along a leading batch axis of the field arrays and updates all of them in one launch, which amortizes kernel launch and graph overhead for small grids. Structures with different meshes or durations are grouped into cohorts of identical shape and executed cohort by cohort, and the results are returned in the input order. A seeded differential-evolution optimizer [24] drives either level through a user-defined objective for gradient-free design.

4 Discrete adjoint

4.1 Transposed updates

Let u^n collect the field, absorber and polarization state after step n and let the forward step be $u^{n+1} = F_n(u^n, p)$ with parameters p . For an objective $J(u^N, \dots)$ the reverse recursion

$$\bar{u}^n = \left(\frac{\partial F_n}{\partial u} \right)^\top \bar{u}^{n+1} + \frac{\partial J}{\partial u^n}, \quad \bar{p} = \sum_n \left(\frac{\partial F_n}{\partial p} \right)^\top \bar{u}^{n+1}, \quad (3)$$

returns the gradient of J with respect to every parameter after one backward sweep. Each transpose is a hand-written kernel that mirrors its forward kernel. The Yee curls transpose to the opposite-direction differences with the same metric, the absorber recursion transposes to $\bar{\psi}_{\text{mem}} = \bar{\psi} + \bar{y}$, $\bar{D} = \kappa^{-1} \bar{y} + c \bar{\psi}_{\text{mem}}$ and $\bar{\psi}_{\text{old}} = b \bar{\psi}_{\text{mem}}$ for the memory variable ψ with coefficients b and c , and the trapezoidal dispersion update transposes to the same two-by-two system applied to the cotangents. The material cotangent $\bar{\epsilon}^{-1}$ accumulates the product of the electric cotangent with the curl at every step, and a chain of differentiable material producers maps it onto design variables: a scalar or tensor permittivity per cell, a density field passed through a physical-radius filter, symmetry averaging and a smooth projection, an analytic solid whose occupancy is a smooth signed-distance step of fixed physical width, a polygon or spline outline whose vertices or control points move that distance, or a source waveform. Objectives may combine point signals, plane spectra, the amplitudes of guided modes at any number of ports, and the far or near-zone fields of a projected box, since each observation is linear in the stored fields and carries its own transpose.

4.2 Checkpoint replay and reversibility

Storing u^n for all n is impossible for three-dimensional problems, so the backward sweep replays the forward pass between checkpoints. With C checkpoints and N steps the binomial schedule of Griewank and Walther [17] minimizes the number of recomputed steps, and the package places the checkpoints on device, host or file storage in an explicit tier order. A retained backward pass keeps the adjoint state between transposes so that several objectives can be differentiated from one forward run. For lossless

nondispersive periodic problems an opt-in reversible mode reconstructs the forward fields by stepping Eq. (1) backward, in the spirit of Mahlau et al. [16], and the absorbing-face variant records the fields at the interface between interior and absorber so that the interior reconstruction stays lossless while the absorber is replayed. Which mode is admissible is decided by the material and boundary content of the project rather than by the user.

5 Tiered-memory execution

5.1 Slabs, halos and banks

The streamed mode partitions the domain along x into slabs of W cells and advances each slab for K steps before moving to the next, so a block of K steps needs a halo of K cells on either side of the slab. Only the slab under update, its halos and its checkpoint reside on the device. The remaining field state lives in banks, contiguous buffers in host memory or in files on solid-state storage, that are read into and written from the device tile by tile. The transposed sweep runs the same tiling in reverse, and both sweeps see the same number of field updates plus the halo overhead $2K/W$ per block. Absorber memories, dispersion states and the stored magnetic-wall faces travel with their slabs, and the tile whose core ends at the upper boundary owns the face arrays.

The memory a run may consume is reserved before any field array is allocated. The reservation counts the file banks that can be alive at once, the host tensors of the material gradient and the device tiles. Both counts are measured. Weak references to every bank, taken with cyclic garbage collection disabled over sixty-eight nondispersive and dispersive layouts, never exceed two banks in the forward sweep and $C + 3$ banks in the backward sweep, and the bound is reached, so the disk reservation charges $C + 3$ states instead of the $C + 5$ it charged before the measurement. An allocation ledger over two recorded runs finds no full-size host tensor in the forward sweep and exactly two in the backward sweep, so a contiguous scalar-permittivity problem reserves four copies of the parameter array rather than eight. Injected allocation, read, write, transfer and reduction failures leave no scratch file behind while their tracebacks are alive. A metadata planner evaluates candidate (W, K, C) policies from the replay work and the logical file traffic they imply before anything is allocated, and can select a policy by an explicit work metric.

5.2 Durable restart

A multi-hour adjoint on a shared workstation should survive the loss of its process. The streamed mode therefore writes a journal at block boundaries: the forward state and partial signals after each forward block, the completed forward signals, and the adjoint state with the partial material gradient after each backward block. A record is written to a temporary directory, synchronized to storage, renamed into place and published through a per-phase pointer file, and the previous record of the same phase is removed afterwards. The journal carries a contract with the hash of the scene, the permittivity, the options and the runtime source files, so a resume with a changed input or a changed solver is refused with the differing keys named. A resumed backward sweep replays only the blocks below the recorded one from the all-zero initial state, so no restart bank has to survive the failure. The design follows the block-granular offloading of large-model inference systems [25, 26], with the causal halo of the wave equation replacing their token dependency.

6 Workbench and interoperability

A validated Project model with a region, materials, structures, sources and monitors is the single description of a scene. Its JSON form is edited in a browser workbench with an objects tree, three orthogonal and one perspective CAD view, a material database, layout and analysis modes, a field viewer




Figure 1: Plane-monitor spectra of the lossless slab against Eq. (4). Absolute errors of the 31 recorded samples are plotted on the logarithmic right axis.

and monitor traces, and the same JSON is loaded by the Python interface, stored with every result and exported as a runnable script. The browser talks to a local FastAPI server that queues jobs and streams their state, and a remote GPU is reached through an SSH tunnel rather than by exposing the server.

GDSII files are imported through the layer stack described in Section 2 and exported from native polygons and rectangles. Port markers in a layout become a runnable two-port or multiport mode network in one call. An independent reader decodes the documented layout records of a commercial project format into the native model and an independent writer splices edited geometry, sources, monitors and region settings back into the original byte stream, reparsing the result before it is returned. These readers and writers cover the documented subset and reject unsupported records.

7 Validation

7.1 Analytic and reference problems

A lossless slab of index 1.5 and thickness $0.2 \mu\text{m}$ in air, excited by a broadband sheet in a periodic two-dimensional cell with a $0.025 \mu\text{m}$ mesh and 16 absorbing cells at each end, is compared with the normal-incidence transmission

$$T(\lambda) = \left[1 + \left(\frac{n^2 - 1}{2n} \right)^2 \sin^2 \left(\frac{2\pi n d}{\lambda} \right) \right]^{-1}, \quad R = 1 - T, \quad (4)$$

at 31 wavelengths between 1.3 and $1.8 \mu\text{m}$. The plane-monitor spectra reach maximum absolute errors of 0.00153444 in transmission and 0.00153535 in reflection, and the residual of $R + T - 1$ stays below 4.43991×10^{-6} (Fig. 1).

A closed total-field/scattered-field box around a dielectric sphere yields the scattering cross-section over nine wavelengths, which is compared with the Mie series [27]. The error over a mesh sequence is not monotonic: it falls from 8.9% at a $0.1 \mu\text{m}$ mesh to 0.31% at $0.05 \mu\text{m}$, rises to 1.05% at $0.025 \mu\text{m}$ and falls again to 0.55% at $0.02 \mu\text{m}$, while doubling the physical duration leaves the $0.025 \mu\text{m}$ value unchanged (Table 1, Fig. 2). The staircased sphere surface, not the box or the duration, sets this error.

Radiation projections are checked against the closed-form fields of an electric dipole. With 14, 28 and 56 samples per face axis the finite-distance projection reaches relative errors of 4.6×10^{-3} , 1.1×10^{-3} and 2.8×10^{-4} at five points between 0.55 and $1.35 \mu\text{m}$ outside the box, a second-order sequence, and the asymptotic and finite-distance models agree to 1.9×10^{-4} at a radius of $2000 \mu\text{m}$. A lossy exterior with index $1.3 + 0.05i$ reproduces the analytic near and far fields with the same convergence. For an open

Table 1: Closed-box TFSF sphere scattering on RTX 5880. The maximum relative cross-section error is taken over nine wavelengths. Physical domain and PML thickness stay fixed. The final row doubles the physical duration. The error sequence is not monotonic.

h (μm)	Grid	Time (fs)	Max error (%)	Empty box (μm^2)
0.1	32^3	120	8.8963	6.68×10^{-15}
0.05	64^3	120	0.3086	1.11×10^{-14}
0.025	128^3	120	1.0474	2.19×10^{-15}
0.02	160^3	120	0.5514	1.37×10^{-14}
0.025	128^3	240	1.0474	2.32×10^{-15}




Figure 2: Scattering cross-section of the dielectric sphere from the closed total-field/scattered-field box at four mesh steps, with the Mie series as reference.

surface the single-face projection of a Gaussian-apodized directional emitter converges to the closed-box result as the field on the omitted faces decays, from 9.6×10^{-2} to 3.8×10^{-3} as the beam waist grows from 0.4 to 0.8 wavelengths, whereas an isotropic dipole stays at 0.79 under refinement.

Guided-mode ports are validated on a straight guide, an interface between guides of index 1.0 and 1.2 whose complex scattering parameters are compared with an independent discrete interface solution, a symmetric Y branch and a four-port crossing. A narrow-aperture port on a wide domain matches the full-cell network to 1.6×10^{-5} in the scattering parameters, the Y branch obeys reciprocity to 2.0×10^{-3} with a 1.2×10^{-4} arm asymmetry, and the crossing obeys reciprocity to 2.4×10^{-4} with a 1.2×10^{-7} gap between its four-fold rotated columns. The power missing from the guided channels of the Y branch is radiated by the staircased junction and is reported as a defect rather than corrected.

The anisotropic path is checked against a birefringent slab of principal indices 1.5 and 2.0 in a co-aligned exterior of indices 1.2 and 1.4 that fills both absorbing faces along the propagation axis. The two polarization channels reproduce the analytic transmission to 1.37% and 0.33%, and the derivative of the transmission with respect to the slab index agrees with the analytic derivative to 6.9%. A magnetic wall next to an absorbing face reflects 1.7×10^{-4} of the incident power with the default profile and 3.5×10^{-4} with $\kappa = 3$, $\alpha = 0.05$ and a quadratic grading.

7.2 Gradient checks

Every differentiable path is compared with full Torch automatic differentiation on a small grid and with central differences along a random direction, and the Taylor remainder of the directional derivative is

Table 2: Recorded gradient agreement of the discrete adjoint. Relative errors compare the transposed sweep with full Torch automatic differentiation unless a finite-difference (FD) reference is named. FP64 on CPU unless stated.

Path	Reference	Relative error
Polygon solid through the checkpointed adjoint	autograd, CUDA FP32	1.2×10^{-7}
Magnetic-wall faces with CPML	autograd	2×10^{-16}
Dispersive faces, four pole parameters	autograd	2.7×10^{-16}
Polygon vertices, voxel VJP	central FD, CPU and CUDA	6.7×10^{-8}
Polygon vertices through the solver	central FD	4.4×10^{-9} to 5.5×10^{-8}
Spline control points	central FD	1.8×10^{-8}
Modal $ S_{21} ^2$ objective, N-port	autograd, CPU and CUDA	1.3×10^{-15} to 5.4×10^{-15}
Far-field intensity objective	central FD	3.4×10^{-10}
Near-zone Poynting objective	central FD	2.7×10^{-9}
Streamed versus resident, tensor media	resident sweep	1×10^{-10}
Streamed versus resident, magnetic walls	resident sweep	3×10^{-7}

required to fall quadratically. Table 2 lists the recorded agreement. The material gradient of the 256^3 problem of Section 8 agrees with the complete-history reference to a relative L_2 error of 2.4×10^{-7} for the dielectric and 4.3×10^{-7} for the two-pole material, both in single precision.

The discrete adjoint of a shape parameter converges to the physical shape derivative only with the mesh. For a two-dimensional pentagon grating of permittivity 4 and period $0.6 \mu\text{m}$ at normal incidence and $1.55 \mu\text{m}$, the vertex gradient at mesh steps of 0.04 , 0.02 and $0.01 \mu\text{m}$ differs from a central-difference reference at $0.005 \mu\text{m}$ by 11.6%, 1.8% and 0.5% for a $0.08 \mu\text{m}$ smoothing width, and by 5.4% and 1.0% at the two coarser steps for a $0.04 \mu\text{m}$ width. Halving the reference step moves the reference by less than 0.07%.

7.3 Test suite

The behaviour above is fixed by 2,141 tests that run on a Linux CPU continuous-integration host and on two CUDA workstations, and every device path is compared with its CPU counterpart. Recorded numerical fixtures carry the SHA-256 of the source files that produced them, so a change in a kernel that alters a recorded number is caught by the test that reads the record.

8 Performance

All timings below were taken on one NVIDIA RTX 5880 Ada with 48 GiB (51.5 GB) of device memory in a Windows workstation with 137 GB of host memory, in single precision unless stated, as medians of three warmed repetitions that include construction, graph capture, stepping and the transfer of the final fields. Cold compilation and disk writes are excluded.

8.1 Single simulations and ensembles

The fused engine is compared with the PyTorch FDTD package it builds on, executed through the same CUDA Graph adapter so that both run without Python in the step loop [19]. Over eight fixtures at 64^3 and 96^3 cells and 800 steps the fused kernels are 15.8 to 17.1 times faster at 64^3 and 9.6 to 12.1 times faster at 96^3 , with point-trace differences below 0.04% that arise from the different absorber stencils (Table 3). Sixteen-case parameter sweeps run 31 to 44 times faster than the sequential upstream package at 32^3 and 15 to 16 times faster at 64^3 when the cases share one tensor batch, and four sphere cases at 64^3 and 800 steps take 1.18 s with one CUDA worker against 47.0 s with one NumPy worker (Fig. 3). The tensor batch

Table 3: Eight forward fixtures on RTX 5880. The external baseline is flaport/fdtd 0.2.2 with an added CUDA Graph adapter. Values are medians of three warmed full solves. The last column is point-trace relative L2 difference in percent, not analytic error. The PML interface stencils differ and final fields are not identical.

Fixture, grid	Upstream graph (ms)	Native fused (ms)	Speedup	Trace L2 (%)
Vacuum, 64^3	649.23	38.06	17.06×	0.0142
Sphere, 64^3	629.07	39.72	15.84×	0.0233
Slab, 64^3	636.26	37.23	17.09×	0.0197
Waveguide, 64^3	632.59	39.90	15.85×	0.0372
Vacuum, 96^3	954.98	78.77	12.12×	0.0120
Sphere, 96^3	848.58	88.05	9.64×	0.0208
Slab, 96^3	836.66	80.37	10.41×	0.0174
Waveguide, 96^3	848.07	80.00	10.60×	0.0328

Table 4: Native sphere ensembles on RTX 5880 with a real CUDA batch axis. Full wall times include setup, transfers and scalar objective evaluation. Final E/H arrays and point traces match independent native runs bitwise. Values below unity in the last column are slowdowns.

Grid, cases	Sequential (ms)	Tensor cohort (ms)	Cases/s	Speedup
32^3 , 1	23.22	20.38	49.06	1.14×
32^3 , 2	46.00	31.39	63.71	1.47×
32^3 , 4	106.81	63.67	62.83	1.68×
32^3 , 8	181.89	88.24	90.66	2.06×
32^3 , 16	341.29	142.46	112.31	2.40×
64^3 , 1	39.08	40.49	24.70	0.97×
64^3 , 2	86.40	78.29	25.55	1.10×
64^3 , 4	167.68	153.87	26.00	1.09×
64^3 , 8	307.35	357.10	22.40	0.86×
64^3 , 16	669.53	852.94	18.76	0.78×

itself is 2.4 times faster than sequential native execution for sixteen 32^3 cases and slower than sequential execution for sixteen 64^3 cases, so cohorts of four are selected at that size (Table 4).

8.2 Other solvers on the same workstation

Meep 1.34 [7] on twelve MPI ranks of an i7-12700 and FDTD 0.6.2 [16] on the RTX 3060 of the same workstation were run, one driver per solver, on the slab of Section 7, a Mie sphere at a $0.05 \mu\text{m}$ mesh, and the vacuum and sphere scenes of Table 3 at 64^3 and 96^3 cells. The absorbing layers share only their thickness, since FDTD grades its convolutional profile differently and Meep uses a stretched-coordinate layer. Only TorchFDTD closes a total-field/scattered-field box around the sphere, so the other two subtract an empty run from total-field planes. Meep runs in double precision, the GPU solvers in single, and the Meep adjoint, a frequency-domain method, was not run. Between any two solvers the slab transmission differs by at most 1.6×10^{-5} and the sphere cross sections by at most 0.7% (Table 5).

On the same device the fused engine steps 4.4 to 5.7 times faster than FDTD and completes a solve 6.4 to 7.3 times faster, and against the twelve Meep ranks, timed with other processes on the host, the full solve is 38 to 44 times faster (Table 6). For a periodic dielectric slab at 64^3 cells and 128 steps the checkpointed sweep returns the permittivity gradient 52 times faster than the FDTD checkpointed




Figure 3: Wall time and throughput of four-structure ensembles. Bars are medians of three measured batches after warm-up and the left axis is logarithmic.

Table 5: TorchFDTD, FDTD_X 0.6.2 and Meep 1.34 on one workstation with an i7-12700, on which Meep ran with twelve MPI ranks, and an RTX 3060 for the two GPU solvers. The slab columns are the maximum absolute errors of the transmission and reflection spectra against Eq. (4) over 31 wavelengths. The sphere column is the maximum relative cross-section error of a sphere of index 2.0 and radius $0.6 \mu\text{m}$ against the Mie series over nine wavelengths at a $0.05 \mu\text{m}$ mesh, with the source and monitor method of each solver. The last row is the largest difference between any two solvers. Meep propagates double-precision fields, the other two single precision.

Solver	Sphere source and monitors	Slab $ \Delta T _{\text{max}}$	Slab $ \Delta R _{\text{max}}$	Sphere error (%)
TorchFDTD	closed TFSF box, scattered-field planes	1.53×10^{-3}	1.54×10^{-3}	1.35
FDTD _X	one-way plane source, empty run subtracted	1.54×10^{-3}	1.53×10^{-3}	1.15
Meep	current sheet, empty run subtracted	1.55×10^{-3}	1.53×10^{-3}	1.15
Largest pairwise difference		1.60×10^{-5}	7.11×10^{-6}	0.68

method with the same two checkpoints and 2.0 times faster than its reversible method, and the gradients agree to a relative L_2 difference of 1.1×10^{-7} .

8.3 Adjoint

A 256^3 problem with 128 steps and two checkpoints, timed from construction through the plane-flux objective and every material cotangent, takes 407.7 s for a dielectric and 742.0 s for a two-pole material on the same Torch model with twelve CPU threads, 2.28 s and 5.00 s resident on the device, and 9.44 s and 23.9 s in the streamed mode with host banks. The resident device adjoint is therefore 179 and 149 times faster than the CPU model, and the streamed mode retains 43 and 31 of those factors while holding only the active slabs on the device.

8.4 Beyond the device memory

Two runs exceed the 51.5 GB of the processor. The first advances a $1152 \times 1024 \times 2048$ real single-precision grid of 2.42 billion cells, whose electric and magnetic state occupies 58.0 GB and 58.3 GB with the absorber memories, for ten steps and a full spatial permittivity gradient with file banks on solid-state storage. The peak device allocation is 2.23 GB, the process peak resident memory is 31.9 GB, the run completes in 58.4 minutes, and the gradient over the causal cone of a point source agrees with a reference

Table 6: Forward throughput and time to gradient on the workstation of Table 5. The forward rows are the sphere scene at 800 steps as medians of three warmed solves, where the full solve includes setup and the final field transfer and the stepping column is the time-stepping loop alone. The adjoint rows are the permittivity gradient of a periodic dielectric slab at 64^3 cells and 128 steps, and the last column is the relative L_2 difference of each gradient from the TorchFDTD gradient. Meep ran in double precision on twelve MPI ranks with other processes present on the host, and its adjoint is excluded because it is a frequency-domain method. The GPU solvers ran in single precision.

Grid	Solver	Full solve (s)	Stepping (s)	Cell-steps/s stepping	Peak device (MiB)
64^3	TorchFDTD	0.114	0.094	2.22×10^9	18
64^3	FDTDx	0.780	0.418	5.02×10^8	89
64^3	Meep	4.373	4.298	4.89×10^7	CPU
96^3	TorchFDTD	0.310	0.270	2.63×10^9	57
96^3	FDTDx	1.998	1.507	4.70×10^8	288
96^3	Meep	13.602	13.356	5.30×10^7	CPU

Adjoint method	Time to gradient (s)	Peak device (MiB)	Gradient relative L_2 difference
TorchFDTD, checkpointed, 2 checkpoints	0.102	40	reference
FDTDx, checkpointed, 2 checkpoints	5.268	303	7.1×10^{-8}
FDTDx, reversible	0.209	111	1.0×10^{-7}

Table 7: Interrupted and resumed single-precision adjoint on the $1152 \times 1024 \times 1920$ grid with slab width 4, temporal depth 4, no interior checkpoints, file banks and one journal record per block.

	Killed process	Resumed process
Forward sweep (s)	559.1	7.7 (restored)
Journal writes (s)	96.6	64.8
Total elapsed (s)	1137.4	957.5
Peak device allocation (GB)		3.03
Peak process resident memory (GB)	28.5	30.0
Backward file traffic, read / written (GB)		546 / 437
Gradient relative L_2 error		9.08×10^{-8}

computed on the cone alone to a relative L_2 error of 9.1×10^{-8} while no cell outside the cone receives a gradient.

The second run tests the journal on a $1152 \times 1024 \times 1920$ grid of 2.26 billion cells with the same policy and one journal record per block (Table 7). The first process exits deliberately after publishing its first backward record. A second process restores the completed forward signals in 7.7 s, resumes the backward sweep below the recorded block and finishes in 957 s with a peak device allocation of 3.03 GB, and its gradient matches the causal-cone reference to 9.08×10^{-8} . Whole-machine counters sampled every five seconds over both processes peak at 38.9 GB in use of 137 GB installed, with a mean disk write rate of 0.59 GB/s and 1.52 TB written in total, so the operating-system file cache stays below 0.3 GB because every bank is written and read once per block.

9 Discussion

We demonstrate a time-domain solver in which the same fused CUDA kernels serve the forward run, its transposed sweep and a streamed execution whose field state exceeds the memory of the processor, with a device footprint of 2 to 3 GB for a 54 GB problem. The measured reservations and the journal turn a large

adjoint from a single fragile allocation into a sequence of block records that can be planned, budgeted and resumed.

The design of the differentiation layer separates it from frameworks that differentiate their own operations. TORCWA and FDTD_X obtain gradients from the framework at the cost of keeping every differentiated operation inside it [15, 16], whereas TorchFDTD pays that cost once, in the transposes of a small set of kernels, and is then free to fuse, tile and offload. On the RTX 3060 comparison of Section 8.2 the same gradient arrived 2.0 times faster than with the reversible FDTD_X method and 52 times faster than with its two-checkpoint method. The price is coverage: a new physical update needs its transpose before it becomes differentiable, and the eigenmode solver and the source positions are held fixed here as they are in the frameworks above. Single-domain multi-GPU execution is implemented on CPU ranks and awaits a validation on two devices.

Several directions follow from the records. The tensor and magnetic-wall paths run their transposes as Torch operations rather than as fused kernels, so their device throughput is bounded by memory traffic that the fused isotropic kernels avoid. The staircase junction of the Y branch and the non-monotonic sphere errors point to the subpixel operator as the next accuracy step for curved and angled interfaces. Sustained throughput of the streamed mode over hours, rather than the ten-step capacity gates recorded here, will decide how far a single workstation can carry an inverse design whose state exceeds its device.

10 Conclusion

TorchFDTD provides an FDTD engine whose forward kernels, transposes and storage tiers are designed together, so that inverse design on one graphics processor is limited by host memory and storage rather than by device memory. The package, its validation records and the source hashes of every measurement are available under the MIT licence [18]. Extending the fused transposes to tensor media and magnetic walls, measuring multi-device execution on real hardware, and running the streamed adjoint through a complete multi-hour optimization are the next steps toward design problems of ten billion cells on a single workstation.

Data and code availability

The source code, the browser workbench, the test suite and every JSON record quoted in this paper are in the public repository [18]. Each record stores the SHA-256 of the source files that produced it and the hardware on which it was taken.

References

- [1] Kane S. Yee. Numerical solution of initial boundary value problems involving Maxwell's equations in isotropic media. *IEEE Transactions on Antennas and Propagation*, 14(3):302–307, 1966. doi: 10.1109/TAP.1966.1138693.
- [2] Allen Taflov and Susan C. Hagness. *Computational Electrodynamics: The Finite-Difference Time-Domain Method*. Artech House, Norwood, MA, 3 edition, 2005.
- [3] Jean-Pierre Bérenger. A perfectly matched layer for the absorption of electromagnetic waves. *Journal of Computational Physics*, 114(2):185–200, 1994. doi: 10.1006/jcph.1994.1159.
- [4] J. Alan Roden and Stephen D. Gedney. Convolution PML (CPML): An efficient FDTD implementation of the CFS-PML for arbitrary media. *Microwave and Optical Technology Letters*, 27(5):334–339, 2000. URL [https://onlinelibrary.wiley.com/doi/10.1002/1098-2760\(200012\)27:5<334::AID-MOP143E3.0.CO;2B2-A](https://onlinelibrary.wiley.com/doi/10.1002/1098-2760(200012)27:5<334::AID-MOP143E3.0.CO;2B2-A).

- [5] Michal Okoniewski, Michal Mrozowski, and Maria A. Stuchly. Simple treatment of multi-term dispersion in FDTD. *IEEE Microwave and Guided Wave Letters*, 7(5):121–123, 1997. doi: 10.1109/75.569723.
- [6] Ardavan Farjadpour, David Roundy, Alejandro Rodriguez, Mihai Ibanescu, Peter Bermel, John D. Joannopoulos, Steven G. Johnson, and Geoffrey W. Burr. Improving accuracy by subpixel smoothing in the finite-difference time domain. *Optics Letters*, 31(20):2972–2974, 2006. doi: 10.1364/OL.31.002972.
- [7] Ardavan F. Oskooi, David Roundy, Mihai Ibanescu, Peter Bermel, John D. Joannopoulos, and Steven G. Johnson. Meep: A flexible free-software package for electromagnetic simulations by the FDTD method. *Computer Physics Communications*, 181(3):687–702, 2010. doi: 10.1016/j.cpc.2009.11.008.
- [8] Ansys. Ansys Lumerical FDTD. Software, 2025. URL <https://www.ansys.com/products/optics/fdtd>. Accessed 21 September 2026.
- [9] Flexcompute. Tidy3D: hardware-accelerated electromagnetic solver. Software, 2025. URL <https://www.flexcompute.com/tidy3d/>. Accessed 21 September 2026.
- [10] Jakob S. Jensen and Ole Sigmund. Topology optimization for nano-photonics. *Laser & Photonics Reviews*, 5(2):308–321, 2011. doi: 10.1002/lpor.201000014.
- [11] Alexander Y. Piggott, Jesse Lu, Konstantinos G. Lagoudakis, Jan Petykiewicz, Thomas M. Babinec, and Jelena Vučković. Inverse design and demonstration of a compact and broadband on-chip wavelength demultiplexer. *Nature Photonics*, 9:374–377, 2015. doi: 10.1038/nphoton.2015.69.
- [12] Sean Molesky, Zin Lin, Alexander Y. Piggott, Weiliang Jin, Jelena Vučković, and Alejandro W. Rodriguez. Inverse design in nanophotonics. *Nature Photonics*, 12:659–670, 2018. doi: 10.1038/s41566-018-0246-9.
- [13] Tyler W. Hughes, Momchil Minkov, Ian A. D. Williamson, and Shanhui Fan. Adjoint method and inverse design for nonlinear nanophotonic devices. *ACS Photonics*, 5(12):4781–4787, 2018. doi: 10.1021/acsp Photonics.8b01522.
- [14] Alec M. Hammond, Ardavan Oskooi, Mo Chen, Zin Lin, Steven G. Johnson, and Stephen E. Ralph. High-performance hybrid time/frequency-domain topology optimization for large-scale photonics inverse design. *Optics Express*, 30(3):4467–4491, 2022. doi: 10.1364/OE.442074.
- [15] Changhyun Kim and Byoungcho Lee. TORCWA: GPU-accelerated Fourier modal method and gradient-based optimization for metasurface design. *Computer Physics Communications*, 282: 108552, 2023. doi: 10.1016/j.cpc.2022.108552.
- [16] Yannik Mahlau, Frederik Schubert, Konrad Bethmann, Reinhard Caspary, Antonio Calà Lesina, Marco Munderloh, Jörn Ostermann, and Bodo Rosenhahn. A flexible framework for large-scale FDTD simulations: open-source inverse design for 3D nanostructures. arXiv:2412.12360, 2024.
- [17] Andreas Griewank and Andrea Walther. Algorithm 799: revolve: an implementation of checkpointing for the reverse or adjoint mode of computational differentiation. *ACM Transactions on Mathematical Software*, 26(1):19–45, 2000. doi: 10.1145/347837.347846.
- [18] Hyoseok Park. TorchFDTD source code and validation records. GitHub repository, MIT licence, 2026. URL <https://github.com/hyoseokp/TorchFDTD>. Accessed 21 September 2026.
- [19] Floris Laporte et al. fdtd: A three-dimensional electromagnetic FDTD simulator in Python, version 0.2.2. MIT-licensed software, 2023. URL <https://github.com/flaport/fdtd>. Accessed 19 September 2026.
- [20] Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An imperative style, high-performance

- deep learning library. In *Advances in Neural Information Processing Systems*, volume 32, pages 8024–8035, 2019.
- [21] Eliane Bécache, Sandrine Fauqueux, and Patrick Joly. Stability of perfectly matched layers, group velocities and anisotropic waves. *Journal of Computational Physics*, 188(2):399–433, 2003. doi: 10.1016/S0021-9991(03)00184-0.
 - [22] Lucas Heitzmann Gabrielli. gdstk: GDSII and OASIS library for Python and C++. Software, 2024. URL <https://github.com/heitzmann/gdstk>. Accessed 21 September 2026.
 - [23] Zhaoming Zhu and Thomas G. Brown. Full-vectorial finite-difference analysis of microstructured optical fibers. *Optics Express*, 10(17):853–864, 2002. doi: 10.1364/OE.10.000853.
 - [24] Rainer Storn and Kenneth Price. Differential evolution: A simple and efficient heuristic for global optimization over continuous spaces. *Journal of Global Optimization*, 11(4):341–359, 1997. doi: 10.1023/A:1008202821328.
 - [25] Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Beidi Chen, Percy Liang, Christopher Ré, Ion Stoica, and Ce Zhang. FlexGen: High-throughput generative inference of large language models with a single GPU. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 31094–31116, 2023.
 - [26] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pages 611–626, 2023. doi: 10.1145/3600006.3613165.
 - [27] Craig F. Bohren and Donald R. Huffman. *Absorption and Scattering of Light by Small Particles*. Wiley, New York, 1983.